

# Problem Set 6

Alexander Brandt

SID: 24092167

October 19 2015

Friendly Collaborators: Milos Atz, Alex Ojala. Most of our discussion centered around RSQLite and Spark syntax. Though I did convince some new folks to try CyberDuck.

## 1

The file takes almost 40 minutes to generate on an xl.large instance, and balloons from 1.7 Gb to about 18 Gb. The code to generate is detailed in the main chunk.

Listing 1: ls results

```
-rw-r--r--  1 ubuntu ubuntu 18G Nov  2 12:14 FlightDatabase.sqlite
```

## 2

### 2.1

The filtering step (detailed in the main code, below) is pretty much the same for both steps: it consists of removing NA's, and in the case of spark, getting rid of header lines from the CSV. Both are easily accomplished with a few lines to remove the offending files. In the case of R, we tag each NA as a numeric code (0.1234) before deleting from the database.

### 2.2

The code for Spark and PySpark is largely modified from what Chris gave us in Unit 7. Instead of computing a "median" for the time delay, we instead bin the times, then map the key/tuple to a string and write it to a file in the hadoop filesystem. By and large, the SPARK method is incredibly faster (all operations can be performed within a few minutes, vs. the arduously long loading times needed by R). For Spark, the output has been slightly modified to highlight the runtimes, without the verbosity.

For SPARK/PySPARK:

```
# Import all our necessary pacakges
import time
from operator import add
import numpy as np
from pyspark import SparkContext

sc = SparkContext()

# Read in all of our .bz2 files
```



```

print "Query start/stop times..."
# Actually create the keys
output = lines.map(computeKeyValue).reduceByKey(add)
# Used for testing:
# print output.collect()[0:10]

# Create the bins based on keys
print "Delay binning start/stop times..."
myResults = output.map(binFun)
# Print to text file
myResults.map(printable).repartition(1).saveAsTextFile('/data/airline_processed')

```

Output:

#### Listing 2: SPARK timing

```

15/11/02 10:15:47 INFO spark.SparkContext: Running Spark version 1.5.1
15/11/02 10:15:48 WARN spark.SparkConf:
SPARK_WORKER_INSTANCES was detected (set to '1').
This is deprecated in Spark 1.0+.

```

Please instead use:

- ./spark-submit with --num-executors to specify the number of executors
- Or `set SPARK_EXECUTOR_INSTANCES`
- `spark.executor.instances` to configure the number of instances in the spark config.

```

15/11/02 10:15:48 INFO spark.SecurityManager: Changing view acls to: root
15/11/02 10:15:48 INFO spark.SecurityManager: Changing modify acls to: root
15/11/02 10:15:48 INFO spark.SecurityManager: SecurityManager: authentication disabled
15/11/02 10:15:49 INFO slf4j.Slf4jLogger: Slf4jLogger started
15/11/02 10:15:49 INFO Remoting: Starting remoting
15/11/02 10:15:49 INFO Remoting: Remoting started; listening on addresses :[akka.tcp://
15/11/02 10:15:49 INFO util.Utils: Successfully started service 'sparkDriver' on port
15/11/02 10:15:49 INFO spark.SparkEnv: Registering MapOutputTracker
15/11/02 10:15:49 INFO spark.SparkEnv: Registering BlockManagerMaster
[...]
15/11/02 10:26:13 INFO remote.RemoteActorRefProvider$RemotingTerminator: Remote daemon

```

And if you wanted to see a few lines from the built file (note, they are in no way ordered):

#### Listing 3: head of SPARK output

```

Alexanders-MBP:airline_processed Alex$ head part-00000
DL,JFK,ORD,10,6,18,2,0,0,24,0.083333
YV,CLT,GSO,1,7,00,1,1,1,1,1.000000
AA,HDN,DFW,3,7,13,2,0,0,74,0.027027
OH,JFK,DTW,5,7,19,1,1,0,9,0.111111
OO,ICT,DEN,2,5,14,1,1,0,1,1.000000
AA,DFW,ATL,9,5,12,4,3,0,90,0.044444
EV,ATL,SHV,10,2,13,0,0,0,3,0.000000
US,SFO,PHL,1,6,22,0,0,0,17,0.000000
HP,LAS,MIA,1,3,22,0,0,0,1,0.000000
AS,SEA,GEG,9,5,23,2,1,0,32,0.062500

```

For R/RSQLite:

```

# Build up the filenames for year.csv.bz2 for our range
years <- seq(from=1987, to=2008)
years_strings <- sapply(years, toString)
fns <- sapply(years_strings, paste, sep="", ".csv.bz2")

# install.packages("RSQLite")
library("RSQLite")
# install.packages("stringr")
library("stringr")
my_path <- "~/ "
setwd(my_path)
# Create our flight database file
database_filename = "FlightDatabase.sqlite"
# Read in based on the filenames
ptm <- proc.time()
for (i in seq(length(fns)))
{
  print(fns[[i]])
  my_bz <- bzfile(fns[[i]])
  my_csv <- read.csv(my_bz, header=TRUE)

  my_csv[is.na(my_csv)] <- 0.1234
  my_csv$DepTime <- substr(str_pad(my_csv$DepTime, 4, pad="0"), 1, 2)

  drv <- dbDriver("SQLite")
  db <- dbConnect(drv, dbname = database_filename)

  dbWriteTable(conn = db, name = "flight_info",
               value = my_csv, row.names = FALSE, append = TRUE)
}
proc.time() - ptm

dbSendQuery(db, "delete from flight_info where DepDelay==0.1234")
dbSendQuery(db, "delete from flight_info where DepTime is '0.'")

# 1)

# I ended up using 'ls' but here is a way to check the file
# size using R itself
file.info(database_filename)

# Query

# 2b) Query based on the departure delays using a sum/case to
# count the number of "offending" flights
ptm <- proc.time()
x <- fetch(dbSendQuery(db, "select UniqueCarrier, Origin,
                          Dest, Month, DayOfWeek,
                          DepTime,
                          SUM(CASE WHEN DepDelay > 30 THEN 1 ELSE 0 END) as DelayedCounts,
                          Count(*) as TotalFlightCounts,
                          CAST(SUM(CASE WHEN DepDelay > 30 THEN 1.0 ELSE 0.0 END) AS FLOAT) / Count(*)
                          as DelayFraction

```

```

        from flight_info group by
        UniqueCarrier, Origin, Dest, Month, DayOfWeek, DepTime
        order by DelayFraction desc"),n=-1)
y <- fetch(dbSendQuery(db, "select UniqueCarrier, Origin,
        Dest, Month, DayOfWeek,
        DepTime,
        SUM(CASE WHEN DepDelay > 60 THEN 1 ELSE 0 END) as DelayedCounts,
        Count(*) as TotalFlightCounts,
        CAST(SUM(CASE WHEN DepDelay > 60 THEN 1.0 ELSE 0.0 END) AS FLOAT) / Count(*)
        as DelayFraction
        from flight_info group by
        UniqueCarrier, Origin, Dest, Month, DayOfWeek, DepTime
        order by DelayFraction desc"),n=-1)

z <- fetch(dbSendQuery(db, "select UniqueCarrier, Origin,
        Dest, Month, DayOfWeek,
        DepTime,
        SUM(CASE WHEN DepDelay > 180 THEN 1 ELSE 0 END) as DelayedCounts,
        Count(*) as TotalFlightCounts,
        CAST(SUM(CASE WHEN DepDelay > 180 THEN 1.0 ELSE 0.0 END) AS FLOAT) / Count(*)
        as DelayFraction
        from flight_info group by
        UniqueCarrier, Origin, Dest, Month, DayOfWeek, DepTime
        order by DelayFraction desc"),n=-1)

proc.time() - ptm

```

Results for 30, 60, 180 minute delays (I had put "90" where I intended to put "30").

Listing 4: Timing for the loading, initial (non-indexed) query

```

> proc.time() - ptm
      user      system elapsed
3041.288      64.808 3124.991

> proc.time() - ptm
      user      system elapsed
1846.000     141.804 2270.214

```

## 2.3

We perform the same calculation, just using python instead of R. It gives the right answer ONLY if the floating point is cast first to Digit. If the code chunk is run "as is" from the problem statement, it sums to 1.

## 2.4

We add a index, which speeds up our calculation precipitously! Note, I ran this in the middle of the night, and I think the process got hung in an odd way, but the user and system time were much faster. When I tested on just single eyar databases on my local system, adding the key always sped up the searches, sometimes by as much as a factor of 2.

```

# Question 2D)
# Add the key based on our search values:
dbSendQuery(db, "create index delay_index on flight_info

```

```

        (UniqueCarrier, Origin, Dest, Month,
         DayOfWeek, DepTime)")
# How we might REMOVE the key, used for testing
# dbSendQuery(db, "drop index delay_index")

# Run the same searches, now much faster
ptm <- proc.time()
x <- fetch(dbSendQuery(db, "select UniqueCarrier, Origin,
                           Dest, Month, DayOfWeek,
                           DepTime,
                           SUM(CASE WHEN DepDelay > 30 THEN 1 ELSE 0 END) as DelayedCounts,
                           Count(*) as TotalFlightCounts,
                           CAST(SUM(CASE WHEN DepDelay > 30 THEN 1.0 ELSE 0.0 END) AS FLOAT) / Count(*)
                           as DelayFraction
                           from flight_info group by
                           UniqueCarrier, Origin, Dest, Month, DayOfWeek, DepTime
                           order by DelayFraction desc"),n=-1)

y <- fetch(dbSendQuery(db, "select UniqueCarrier, Origin,
                           Dest, Month, DayOfWeek,
                           DepTime,
                           SUM(CASE WHEN DepDelay > 60 THEN 1 ELSE 0 END)
                           as DelayedCounts,
                           Count(*) as TotalFlightCounts,
                           CAST(SUM(CASE WHEN DepDelay > 60 THEN 1.0 ELSE 0.0 END) AS FLOAT) / Count(*) as
                           from flight_info group by
                           UniqueCarrier, Origin, Dest, Month, DayOfWeek, DepTime
                           order by DelayFraction desc"),n=-1)

z <- fetch(dbSendQuery(db, "select UniqueCarrier, Origin,
                           Dest, Month, DayOfWeek,
                           DepTime,
                           SUM(CASE WHEN DepDelay > 180 THEN 1 ELSE 0 END) as DelayedCounts,
                           Count(*) as TotalFlightCounts,
                           CAST(SUM(CASE WHEN DepDelay > 180 THEN 1.0 ELSE 0.0 END) AS FLOAT) / Count(*)
                           as DelayFraction
                           from flight_info group by
                           UniqueCarrier, Origin, Dest, Month, DayOfWeek, DepTime
                           order by DelayFraction desc"),n=-1)

proc.time() - ptm

```

Listing 5: Timing for SQLite query

```

> proc.time() - ptm
      user      system    elapsed
 832.760    139.484 12617.140

```

## 2.5

Using R, we take our object generated by RSQLite and then subset based on flights with at least 150 entries. Then we view the top 10 for each of the 30, 90, 180 minute delays, respectively.

```
# Question 2E)
#
xb <- subset(x, TotalFlightCounts > 149)
yb <- subset(y, TotalFlightCounts > 149)
zb <- subset(z, TotalFlightCounts > 149)

head(xb,n=10)
head(yb,n=10)
head(zb,n=10)
```

Results for 30, 60, 180 minute delays (I had put "90" where I intended to put "30").

Listing 6: Flight Delay (Top 10)

```
head(xb,n=10)
```

	UniqueCarrier	Origin	Dest	Month	DayOfWeek	DepTime	DelayedCounts
1945772	WN	HOU	DAL	2	5	19	61
1946374	WN	DAL	HOU	6	5	20	62
1971173	WN	DAL	HOU	2	5	21	63
1974391	WN	DAL	HOU	5	5	21	61
1985521	WN	HOU	DAL	2	5	20	58
1994481	WN	HOU	DAL	10	5	20	61
1995486	UA	LAX	SFO	12	5	18	52
1997129	WN	DAL	HOU	12	5	20	53
1997219	WN	HOU	DAL	6	5	21	56
1997436	WN	HOU	DAL	6	5	20	56

  

	TotalFlightCounts	DelayFraction
1945772	153	0.3986928
1946374	158	0.3924051
1971173	168	0.3750000
1974391	165	0.3696970
1985521	162	0.3580247
1994481	175	0.3485714
1995486	150	0.3466667
1997129	155	0.3419355
1997219	164	0.3414634
1997436	165	0.3393939

  

```
> head(yb,n=10)
```

	UniqueCarrier	Origin	Dest	Month	DayOfWeek	DepTime	DelayedCounts
1638876	WN	HOU	DAL	6	5	18	36
1666191	WN	HOU	DAL	5	4	21	31
1666878	WN	HOU	DAL	2	5	19	26
1666978	WN	HOU	DAL	10	5	18	33
1732659	WN	HOU	DAL	5	4	19	29
1744973	WN	HOU	DAL	10	5	20	28
1745327	WN	HOU	DAL	6	4	19	28
1749951	WN	DAL	HOU	2	5	21	26
1761811	WN	DAL	HOU	4	5	21	25
1761812	UA	LAX	SFO	10	5	12	23

  

	TotalFlightCounts	DelayFraction
1638876	189	0.1904762
1666191	180	0.1722222
1666878	153	0.1699346
1666978	195	0.1692308

```

1732659          174      0.1666667
1744973          175      0.1600000
1745327          177      0.1581921
1749951          168      0.1547619
1761811          163      0.1533742
1761812          150      0.1533333
> head(zb,n=10)
      UniqueCarrier Origin Dest Month DayOfWeek DepTime DelayedCounts
378918           WN   HOU  DAL     7           7      19             5
383602           WN   HOU  DAL     4           5      20             5
397917           WN   HOU  DAL     4           2      21             4
399799           WN   HOU  DAL     7           3      20             4
403164           WN   DAL  HOU     5           4      19             4
403202           WN   HOU  DAL    10           5      20             4
413930           WN   DAL  HOU     6           2      21             3
414237           AA   ORD  DFW    12           4      18             3
415160           UA   SFO  LAX    10           7      16             3
415161           UA   SFO  LAX    12           7      16             3
      TotalFlightCounts DelayFraction
378918             157      0.03184713
383602             167      0.02994012
397917             161      0.02484472
399799             166      0.02409639
403164             173      0.02312139
403202             175      0.02285714
413930             150      0.02000000
414237             153      0.01960784
415160             153      0.01960784
415161             153      0.01960784

```

### 3

```

# Question 3)
# install.packages("parallel")
library("parallel")

# Build a function, getDelays, which takes a time value
# and a "string" which represents a letter of the alphabet
# this is used in a regex way to break down the search by
# flight code starting letter (handy way to ensure no doubled
# results)
getDelays <- function(x,s) {
  # print(x)
  s <- toString(s)
  # print(s)
  drv <- dbDriver("SQLite")
  db <- dbConnect(drv, dbname = "Big_v4.sqlite")
  query <- sprintf("select UniqueCarrier, Origin,
                    Dest, Month, DayOfWeek,
                    DepTime,
                    SUM(CASE WHEN DepDelay > %i THEN 1 ELSE 0 END) as DelayedCounts,
                    Count(*) as TotalFlightCounts,

```



```

        CAST(SUM(CASE WHEN DepDelay > %i THEN 1.0 ELSE 0.0 END) AS FLOAT) / Count(*)
        as DelayFraction
    from flight_info where Origin like '%s%' group by
    UniqueCarrier, Origin, Dest, Month, DayOfWeek, DepTime
    order by DelayFraction desc", x, x, s)

tmp <- fetch(dbSendQuery(db, query),n=-1)
return(tmp)
}

alphabet = c("A","B","C","D","E","F","G","H","I","J","K","L",
            "M","N","O","P","Q","R","S","T","U","V","W","X",
            "Y","Z")
times <- c(30, 60, 180)
# All combinations of alphabet letters and delay times
tlc <- expand.grid(times,alphabet)
names(tlc) <- c("x","s")

ptm <- proc.time()
question_3 <- mcmapply(getDelays, tlc$x, tlc$s, mc.cores=4)
proc.time() - ptm

```

Listing 7: Timing for all\_preprocess.sh

```

> proc.time() - ptm
      user   system elapsed
245.192    22.508   375.559

```

## 4

The preprocessing I used here was basically derived from my solution in ps2, just generalized to work with more than one bzip2 file. The code and a wrapper script is shown below. The file almost takes about 10 minutes, which is m3.xlarge instance. It is probably worth it, for the case of R, because reading in the files can take a very long time.

Listing 8: all\_preprocess.sh

```

myyear=$1
# Extract the header so we can find our columns of interest
bzip2cat $myyear.csv.bz2 | head -n 1 > $myyear.header.txt
# We will use the file line coordinates as the proxy for index columns
sed -e '$s/,\|\\n/g' $myyear.header.txt > $myyear.header.nsv
# Our desired headers
for i in "UniqueCarrier" "Origin" "Dest" "Month" "DayOfWeek" "DepTime" "DepDelay"
do
    x=`grep -n ^$i$ $myyear.header.nsv | cut -d':' -f 1`
    v="$v $x"
done
echo $v

# Now $v contains our columns of interest, which we just need
# to separate by commas to use with cut. A sed command will
# accomplish this with ease.

```

```
bzcat $myyear.csv.bz2 | \
    cut -d, -f`echo $v | \
    sed 's/ /,/g'` | bzip2 > $myyear.pp.csv.bz2
```

This script is called with:

Listing 9: preprocess.sh

```
date
for f in `seq 1987 1 2008`
do
    echo $f
    ~/preprocess.sh $f
done
date
```

And the result:

Listing 10: Timing for all\_preprocess.sh

```
Mon Nov  2 08:25:32 UTC 2015
Mon Nov  2 08:35:15 UTC 2015
```